

Project 2: Understanding Cache Memories

519021910095, Qianyun Guo, guoqianyun@sjtu.edu.cn

June 2, 2021

1 Introduction

In this project, I had a deep understanding of the working mechanism of cache. In part A, I wrote a cache simulator using C language, which simulates the behavior of a cache memory and records the counts of hit, miss and eviction during the test work. In part B, I tried to optimize a matrix transpose function with the goal of making it cache friendly, which means reduce the miss count during the execution of matrix transpose according to the understanding of cache operation mechanism.

2 Experiments

2.1 Part A

2.1.1 Analysis

This part is about writing a cache simulator using C language. To put it simply, the C program should have the following functions: analyzing the memory trace in the reference trace files as input, maintaining the operation of the cache based on the input, using variables to record the counts of hit, miss, and eviction. The difficult points and code techniques for this part are as follows.

Difficult points

- Analyzing the input from the reference trace files.
- Having a deep understanding of the working mechanism of cache (set-associate).
- Using proper methods to maintaining the operation of the cache.
- The implementation of LRU replacement strategy.

Code techniques

- Set the counts of hit, miss and eviction as global static variables for functions to maintain.
- Design a BLOCK struct to simulate cache block with valid-invalid bit, tag information and LRU record.
- Design a CACHE struct to simulate cache with the number of ways per set, number of sets in the cache, the size of each block and a pointer pointed to the cache address.

- Design a function (str_int) to change the decimal data in a string to an integer.
- Design a function (hex_dec) to change the hexadecimal data in a string to long type, which is used for analyzing memory address.
- Design a function (parse) to analyze the memory trace in the reference trace files and get the operation type and memory address.
- Design a function (execute) to simulate accessing an address in the cache, with the use of LRU replacement strategy. Pay attention to maintaining hit, miss, and eviction variables.
- In function main(), we first analyze the command to get important parameters for initializing the cache, then open the trace files and simulate cache operation using the above functions.
- LRU replacement strategy implementation.
Using the count to record the cumulative memory trace. Each time the program read in a memory trace, update the count. In the function execute(), the count is used to update the LRU record of accessed block. In this case, when cache have to search an eviction, the block that has the minimum LRU record means it is the one the longest not visited, and it would be chosen to be the eviction.

To enhance code readability, I wrote necessary comments for critical steps in my codes, please refer to the code section.

2.1.2 Code

csim.c

```

1 //Guo Qianyun 519021910095
2 #include "cachelab.h"
3 #include <getopt.h>
4 #include <stdlib.h>
5 #include <unistd.h>
6 #include <stdio.h>
7 #include <string.h>
8 #define BUFFERSIZE 50
9 static int hit = 0; //hit count
10 static int miss = 0; //miss count
11 static int eviction = 0; //eviction count
12 struct BLOCK {
13     int valid; //valid-invalid bit
14     long tag; //tag infomation
15     int LRU; //LRU record
16 };
17
18 struct CACHE {
19     int waynum; //E associativity: num of ways per set
20     int setnum; //2^s num of sets
21     int blocksize; //2^b b: num of block bits
22     struct BLOCK **c;
23 };
24 int str_int(char *str); //change decimal str to int
25 long hex_dec(char *str); //change hex str to decimal
26 //analyze the memory trace and get the operation type and memory address

```

```

27 void parse(char *buffer, char *op, long *add);
28 //simulate accessing an address in the cache, using LRU
29 void execute(struct CACHE *cache, long address, int cnt);
30
31
32 int main(int argc, char *argv[])
33 {
34     struct CACHE cache;
35     int s, b, S, E, B = 0;
36     FILE *fp = NULL;           //for tracefile
37     char buffer[BUFFERSIZE];   //for instruction in tracefile
38
39     //get S, E, B
40     for(int i = 0; i < argc; i++)
41     {
42         if(argv[i][0] == '-')
43         {
44             if(argv[i][1] == 's')
45             {
46                 i++;
47                 s = str_int(argv[i]);
48                 S = 1 << s; //2^s
49             }
50             if(argv[i][1] == 'E')
51             {
52                 i++;
53                 E = str_int(argv[i]);
54             }
55             if(argv[i][1] == 'b')
56             {
57                 i++;
58                 b = str_int(argv[i]);
59                 B = 1 << b; //2^b
60             }
61             if(argv[i][1] == 't')
62             {
63                 i++;
64                 if((fp = fopen(argv[i], "r")) == NULL)
65                 {
66                     printf(" ERROR: FILE %s OPEN FAILED", argv[i]);
67                     exit(1);
68                 }
69             }
70         }
71     }
72     if (s <= 0 || E <= 0 || b <= 0)
73     {
74         printf(" ERROR: INVALID PARAMETER");
75         exit(1);
76     }
77
78     //initialize cache
79     cache.waynum = E;
80     cache.setnum = S;
81     cache.blocksize = B;
82     cache.c = (struct BLOCK **) malloc (sizeof(struct BLOCK *) * S); //S sets
83     for (int i = 0; i < S; i++)
84     {

```

```

85     // E ways per set
86     cache.c[i] = (struct BLOCK *) malloc (sizeof(struct BLOCK) * E);
87     //initialize each block
88     for(int j = 0; j < E; j++)
89     {
90         cache.c[i][j].valid = 0;
91         cache.c[i][j].tag = 0;
92         cache.c[i][j].LRU = 0;
93     }
94 }
95
96 int cnt = 0;
97 while(fgets(buffer,sizeof(buffer), fp)) //read in next memory trace
98 {
99     cnt++;
100    char op;
101    long address;
102    parse(buffer, &op, &address); //get operation and address
103    if(op == 'I') continue; //skip instruction cache accesses
104    execute(&cache, address, cnt);
105    if(op == 'M') //M execute twice
106        execute(&cache, address, cnt);
107 }
108 printSummary(hit, miss, eviction);
109 return 0;
110 }
111
112 //change decimal str to int
113 int str_int(char *str)
114 {
115     int len = strlen(str);
116     int res = 0;
117     for(int i = 0; i < len; i++)
118     {
119         res = res * 10 + str[i] - '0';
120     }
121     return res;
122 }
123
124 //change hex str to decimal
125 long hex_dec(char *str)
126 {
127     int len = strlen(str);
128     long res = 0;
129     for(int i = 0; i < len; i++)
130     {
131         if(str[i] >='0' && str[i] <= '9')
132             res = res * 16 + str[i] - '0';
133         if(str[i] >='a' && str[i] <= 'f')
134             res = res * 16 + str[i] - 'a' + 10;
135         if(str[i] >='A' && str[i] <= 'F')
136             res = res * 16 + str[i] - 'A' + 10;
137     }
138     return res;
139 }
140 }
141
142 //simulate accessing an address in the cache, using LRU

```

```

143 void execute(struct CACHE *cache, long address, int cnt)
144 {
145     int set_id = (address / cache->blocksize) % (cache->setnum); //which set
146     long tag_num = (address / cache->blocksize) / (cache->setnum); //tag info
147     //search in the set
148     int pos = -1;
149     for (int i = 0; i < cache->waynum; i++)
150     {
151         //found
152         if (cache->c[set_id][i].valid == 1 && cache->c[set_id][i].tag == tag_num)
153         {
154             hit++;
155             cache->c[set_id][i].LRU = cnt; //update LRU record
156             return;
157         }
158         //not found search if empty
159         if(cache->c[set_id][i].valid == 0)
160             pos = i;
161     }
162     miss++;
163     if(pos >= 0 && pos < cache->waynum)//miss but still have space
164     {
165         cache->c[set_id][pos].valid = 1;
166         cache->c[set_id][pos].tag = tag_num;
167         cache->c[set_id][pos].LRU = cnt; //update LRU record
168         return;
169     }
170     else//evict
171     {
172         eviction++;
173         int min = cache->c[set_id][0].LRU;
174         pos = 0;
175         for(int i = 1; i < cache->waynum; i++) //find eviction (minimum LRU)
176         {
177             if(cache->c[set_id][i].LRU < min)
178             {
179                 min = cache->c[set_id][i].LRU;
180                 pos = i;
181             }
182         }
183         cache->c[set_id][pos].tag = tag_num;
184         cache->c[set_id][pos].LRU = cnt; //update LRU record
185         cache->c[set_id][pos].valid = 1;
186     }
187     return;
188 }
189
190 //analyze the memory trace and get the operation type and memory address
191 void parse(char *buffer, char *op, long *address)
192 {
193     char addr[50];
194     int i = 0;
195     while (buffer[i] == ' ') i++;
196     *op=buffer[i];
197     i++;
198     while (buffer[i] == ' ') i++;
199     int j = 0;
200     for(; buffer[i] != ',';i++,j++)

```

```

201 {
202     addr[j] = buffer[i];
203 }
204 addr[j] = 0;
205 *address = hex_dec(addr); //change string to long type
206 }

```

2.1.3 Evaluation

- Test for `csim.c` (Figure 1)
Test command are as follows.

```

1 make
2 ./test-csim

```

The command tests the correctness of the cache simulator on the reference traces with different cache parameters. As shown in Figure 1, the data of my cache simulator is consistent with the data from reference simulator and get full marks in all the 8 test cases.

```

gqy@gqy-VirtualBox:~/archlab2$ ./test-csim

```

Points (s,E,b)	Your simulator			Reference simulator			
	Hits	Misses	Evicts	Hits	Misses	Evicts	
3 (1,1,1)	9	8	6	9	8	6	traces/yi2.trace
3 (4,2,4)	4	5	2	4	5	2	traces/yi.trace
3 (2,1,4)	2	3	1	2	3	1	traces/dave.trace
3 (2,1,3)	167	71	67	167	71	67	traces/trans.trace
3 (2,2,3)	201	37	29	201	37	29	traces/trans.trace
3 (2,4,3)	212	26	10	212	26	10	traces/trans.trace
3 (5,1,5)	231	7	0	231	7	0	traces/trans.trace
6 (5,1,5)	265189	21775	21743	265189	21775	21743	traces/long.trace
27							

```

TEST_CSIM_RESULTS=27

```

Figure 1: Part A `csim.c`

2.2 Part B

2.2.1 Analysis

This part is about optimizing a matrix transpose function in order to get better cache performance, that is to say, the function causes as few cache-misses as possible. The tests evaluate 3 different-sized matrices: 32*32, 64*64, 61*67 with cache parameters $s = 5$, $E = 1$, $b = 5$. Before starting, it is important to analyze the cache situation according to different sizes. The difficult points and code techniques as well as thoughts for this part are as follows.

Difficult points

- Analyzing the cache situation for each test size with parameters provided.
- Designing proper method to reduce misses.

- Allocate variables reasonably.

Code techniques and thoughts (Classified discussion)

- 32*32
 - Create 8 temporary variables to help transpose the matrices.
 - Block matrix, with each part size 8*8. The reasons why I choose 8*8 are as follows.
 - Since the cache parameters $s = 5$, $E = 1$, $b = 5$. The cache has 32 sets and each set has one block. The size of each block is 32 so each block contains 8 integers.
 - 32*32 matrix, with 32 integers in each line. Since each block contains 8 integers and the cache can contain 32 blocks, the matrix element address is continuous so at most continuous 8 lines in A (8*32 integers in total) can be in the cache at the same time. As shown in Figure 2.

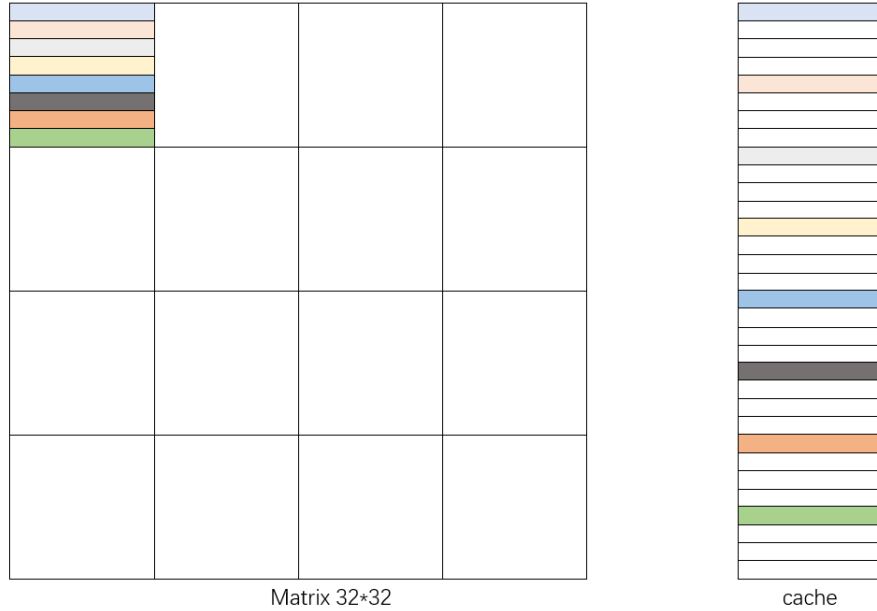


Figure 2: cache situation for 32*32 matrix

- When executing matrix transpose, the operation in B has the same rules as A. That means, to avoid conflict as much as possible, the size of block matrix is at most 8*8. On the other side, if the size of block size is smaller than 8*8, part of the cache space will be wasted, and part of the data in one cached block will be wasted.
- While transpose block matrix, we use the 8 temporary variables instead of using 8 loops like function below for that alternately access matrices A and B may cause unnecessary conflict misses.

```

1  for (i = 0; i < N; i+=8)
2  {
3      for (j = 0; j < M; j+=8)
4      {
5          for (a = i; a < i+8; a++)
6          {
7              for (b = j; b < j+8; b++)
8              {
9                  B[b][a] = A[a][b];
10             }
11         }
12     }
13 }

```

In this way, the total miss count is 287.

- 64*64
 - Create 8 temporary variables to help transpose the matrices.
 - Block matrix. Similar to the above analysis, but this time at most continuous 4 lines in A (4*64 integers in total) can be in the cache at the same time. Thus, the first reaction is to choose the size of 4*4. However, the implementation in this way will cause 1699 miss count in total, which has not yet reached the optimization requirements, so further optimization is needed. As shown in Figure 3.




Figure 3: cache situation for 64*64 matrix

- It can be noticed that when transposing the first 4×4 block matrix, some of the data from A in the cached block has not used yet, and some of the space of B in the cached block has not used either. To make full use of them, we temporarily transpose the unused part of A to unused part of B. In this way, we have to expand the operation range of each loop to 8×8 .
- The specific steps to transpose 8×8 block matrix are as follows.
 - Firstly, transpose the first 4×4 block matrix in A to B (up-left part in A to up-left part in B), and temporarily transpose the up-right 4×4 part in A to the up-right 4×4 part in B for these parts has been cached. As shown in Figure 4.




Figure 4: First Step

- Secondly, transpose the down-left 4×4 part in A to the up-right 4×4 part in B and move the temporary data in the up-right 4×4 part in B to the down-left 4×4 part in B. We have to first use 8 temporary variables to record the data for it may be covered during execution. As shown in Figure 5.




Figure 5: Second Step

- Thirdly, transpose the down-right 4×4 part in A to the down-right 4×4 part in B. As shown in Figure 6.




Figure 6: Third Step

In this way, the total miss count is 1179.

- 61*67
 - We still choose to operate on block matrix, but since M is not equal to N , we can't use fixed 8 variables as before so we chose to use loops.
 - First, I tried 8*8 block matrix and the miss count reached 2118, which has not yet reached the optimization requirements.
 - By accident, I tired 16*16 block matrix and the miss count reached 1992, which has reached the optimization requirements.
 - Noticed that the miss count may have correlation to the size of block matrix size. I defined PARTSIZE and change it to test the correlation. The results are shown in the bellow Table 1

Table 1: the correlation between PARTSIZE and MISS

PARTSIZE	4	8	12	16	17	18	19	20
MISS	2425	2118	2057	1992	1950	1961	1979	2002

- We can see the correlation between miss count and PARTSIZE is a U shape, and when PARTSIZE is 17, we get the lowest point. So finally the PARTSIZE is set to 17.

To enhance code readability, I wrote necessary comments for critical steps in my codes, please refer to the code section.

2.2.2 Code

trans.c

```
1 //Guo Qianyun 519021910095
2 /*
3  * trans.c - Matrix transpose B = A^T
4  *
5  * Each transpose function must have a prototype of the form:
6  * void trans(int M, int N, int A[N][M], int B[M][N]);
7  *
8  * A transpose function is evaluated by counting the number of misses
9  * on a 1KB direct mapped cache with a block size of 32 bytes.
10 */
11 #include <stdio.h>
12 #include "cachelab.h"
13 #define PARTSIZE 17
14 int is_transpose(int M, int N, int A[N][M], int B[M][N]);
15
16 /*
17  * transpose_submit - This is the solution transpose function that you
18  * will be graded on for Part B of the assignment. Do not change
19  * the description string "Transpose submission", as the driver
20  * searches for that string to identify the transpose function to
21  * be graded.
22  */
23 char transpose_submit_desc[] = "Transpose submission";
24 void transpose_submit(int M, int N, int A[N][M], int B[M][N])
25 {
26     if(M == 32 && N == 32)
27     {
28         int t0, t1, t2, t3, t4, t5, t6, t7;
29         for(int i = 0; i < N; i += 8)
30         {
31             for(int j = 0; j < M; j += 8)
32             {
33                 for(int a = i; a < i+8; a++)
34                 {
35                     int b = j;
36                     // 8 integer from a line in A
37                     t0 = A[a][b];
38                     t1 = A[a][b+1];
39                     t2 = A[a][b+2];
40                     t3 = A[a][b+3];
41                     t4 = A[a][b+4];
42                     t5 = A[a][b+5];
43                     t6 = A[a][b+6];
44                     t7 = A[a][b+7];
45                     //put to B
46                     B[b][a] = t0;
47                     B[b+1][a] = t1;
48                     B[b+2][a] = t2;
49                     B[b+3][a] = t3;
50                     B[b+4][a] = t4;
51                     B[b+5][a] = t5;
52                     B[b+6][a] = t6;
53                     B[b+7][a] = t7;
54                 }

```

```

55     }
56 }
57 }
58 else if(M == 64 && N == 64)
59 {
60     int t0, t1, t2, t3, t4, t5, t6, t7;
61     for(int i = 0; i < N; i += 8)
62     {
63         for(int j = 0; j < M; j +=8)
64         {
65             for(int a = i; a < i+4; a++)
66             {
67                 int b = j;
68
69                 t0 = A[a][b];
70                 t1 = A[a][b+1];
71                 t2 = A[a][b+2];
72                 t3 = A[a][b+3];
73                 t4 = A[a][b+4];
74                 t5 = A[a][b+5];
75                 t6 = A[a][b+6];
76                 t7 = A[a][b+7];
77
78                 //up_left part of A to up_left part of B
79                 B[b][a] = t0;
80                 B[b+1][a] = t1;
81                 B[b+2][a] = t2;
82                 B[b+3][a] = t3;
83
84                 //temporarily up_right part of A to up_right part of B
85                 B[b][a+4] = t4;
86                 B[b+1][a+4] = t5;
87                 B[b+2][a+4] = t6;
88                 B[b+3][a+4] = t7;
89             }
90             for(int b = j; b < j+4; b++)
91             {
92                 int a = i;
93                 t0 = A[a+4][b];
94                 t1 = A[a+5][b];
95                 t2 = A[a+6][b];
96                 t3 = A[a+7][b];
97
98                 t4 = B[b][a+4];
99                 t5 = B[b][a+5];
100                t6 = B[b][a+6];
101                t7 = B[b][a+7];
102
103                //down_left part of A to up_right part of B
104                B[b][a+4] = t0;
105                B[b][a+5] = t1;
106                B[b][a+6] = t2;
107                B[b][a+7] = t3;
108
109                //the temporary up_right part of B move to down_left part of B
110                B[b+4][a] = t4;
111                B[b+4][a+1] = t5;
112                B[b+4][a+2] = t6;

```

```

113         B[b+4][a+3] = t7;
114     }
115     for(int a = i+4; a<i+8; a++)
116     {
117         int b = j+4;
118         t0 = A[a][b];
119         t1 = A[a][b+1];
120         t2 = A[a][b+2];
121         t3 = A[a][b+3];
122
123         //down_right part of A to down_right part of B
124         B[b][a] = t0;
125         B[b+1][a] = t1;
126         B[b+2][a] = t2;
127         B[b+3][a] = t3;
128     }
129 }
130 }
131 }
132 else if(M == 61 && N == 67)
133 {
134     for(int i = 0; i < N; i+=PARTSIZE)
135     {
136         for(int j = 0; j < M; j+=PARTSIZE)
137         {
138             for(int a = i; a < i+PARTSIZE && a < N; a++)
139             {
140                 for(int b = j; b < j+PARTSIZE && b < M; b++)
141                 {
142                     B[b][a] = A[a][b];
143                 }
144             }
145         }
146     }
147 }
148 else
149 {
150     int i, j, tmp;
151     for (i = 0; i < N; i++) {
152         for (j = 0; j < M; j++) {
153             tmp = A[i][j];
154             B[j][i] = tmp;
155         }
156     }
157 }
158 }
159
160 /*
161  * You can define additional transpose functions below. We've defined
162  * a simple one below to help you get started.
163  */
164
165 /*
166  * trans - A simple baseline transpose function, not optimized for the cache.
167  */
168 char trans_desc[] = "Simple row-wise scan transpose";
169 void trans(int M, int N, int A[N][M], int B[M][N])
170 {

```

```

171     int i, j, tmp;
172
173     for (i = 0; i < N; i++) {
174         for (j = 0; j < M; j++) {
175             tmp = A[i][j];
176             B[j][i] = tmp;
177         }
178     }
179
180 }
181
182 /*
183  * registerFunctions - This function registers your transpose
184  *                     functions with the driver. At runtime, the driver will
185  *                     evaluate each of the registered functions and summarize their
186  *                     performance. This is a handy way to experiment with different
187  *                     transpose strategies.
188  */
189 void registerFunctions()
190 {
191     /* Register your solution function */
192     registerTransFunction(transpose_submit, transpose_submit_desc);
193
194     /* Register any additional transpose functions */
195     registerTransFunction(trans, trans_desc);
196 }
197
198 /*
199  * is_transpose - This helper function checks if B is the transpose of
200  *               A. You can check the correctness of your transpose by calling
201  *               it before returning from the transpose function.
202  */
203
204 int is_transpose(int M, int N, int A[N][M], int B[M][N])
205 {
206     int i, j;
207
208     for (i = 0; i < N; i++) {
209         for (j = 0; j < M; ++j) {
210             if (A[i][j] != B[j][i]) {
211                 return 0;
212             }
213         }
214     }
215     return 1;
216 }

```

2.2.3 Evaluation

Test command are as follows.

```

1 make
2 ./test-trans -M 32 -N 32
3 ./test-trans -M 64 -N 64
4 ./test-trans -M 61 -N 67

```

```

gqy@gqy-VirtualBox:~/archlab2$ ./test-trans -M 32 -N 32

Function 0 (2 total)
Step 1: Validating and generating memory traces
Step 2: Evaluating performance (s=5, E=1, b=5)
func 0 (Transpose submission): hits:1766, misses:287, evictions:255

Function 1 (2 total)
Step 1: Validating and generating memory traces
Step 2: Evaluating performance (s=5, E=1, b=5)
func 1 (Simple row-wise scan transpose): hits:870, misses:1183, evictions:1151

Summary for official submission (func 0): correctness=1 misses=287
TEST_TRANS_RESULTS=1:287

```

Figure 7: Part B: M=32, N=32

```

gqy@gqy-VirtualBox:~/archlab2$ ./test-trans -M 64 -N 64

Function 0 (2 total)
Step 1: Validating and generating memory traces
Step 2: Evaluating performance (s=5, E=1, b=5)
func 0 (Transpose submission): hits:9066, misses:1179, evictions:1147

Function 1 (2 total)
Step 1: Validating and generating memory traces
Step 2: Evaluating performance (s=5, E=1, b=5)
func 1 (Simple row-wise scan transpose): hits:3474, misses:4723, evictions:4691

Summary for official submission (func 0): correctness=1 misses=1179
TEST_TRANS_RESULTS=1:1179

```

Figure 8: Part B: M=64, N=64

```

gqy@gqy-VirtualBox:~/archlab2$ ./test-trans -M 61 -N 67

Function 0 (2 total)
Step 1: Validating and generating memory traces
Step 2: Evaluating performance (s=5, E=1, b=5)
func 0 (Transpose submission): hits:6229, misses:1950, evictions:1918

Function 1 (2 total)
Step 1: Validating and generating memory traces
Step 2: Evaluating performance (s=5, E=1, b=5)
func 1 (Simple row-wise scan transpose): hits:3756, misses:4423, evictions:4391

Summary for official submission (func 0): correctness=1 misses=1950
TEST_TRANS_RESULTS=1:1950

```

Figure 9: Part B: M=61, N=67

- Test for M=32 and N=32 (Figure 7)
- Test for M=64 and N=64 (Figure 8)
- Test for M=61 and N=67 (Figure 9)

The command tests the count of miss of each type of matrix. The count of miss for M=32 and N=32 is 287, the count of miss for M=64 and N=64 is 1179, the count of miss for M=61 and N=67 is 1950, which all reached the optimization requirements.

3 Conclusion

3.1 Problems

- Have a thorough understanding of the mechanism of cache operation. Writing a cache simulator entails a good grasp of the cache mechanism. Since the `csim.c` is almost empty, we have to write from scratch. At first, I had no clue about where to start, so I chose to review the content of the classes about cache and then figure out the way.
- Figure out implementation method for cache simulation and LRU replacement strategy. To simulate the cache properly we have to create proper functions and structs. Fortunately, I have tried to write a memory simulator with TLB and page table in the operating system class and this part is similar to that, so it can be solved after careful thoughts.
- Figure out further optimization method for 64*64 matrix transpose. When noticing that 4*4 block matrix method is not enough for the optimization requirements, I was stuck in a bottleneck. After an efficient discussion with classmates and a deep personal thought, the further optimization method was figured out finally.

3.2 Achievements

- In Part A, I successfully wrote a cache simulator which runs correctly with right record of hit count, miss count and eviction count, using LRU replacement strategy.
- In Part B, I carefully analyzed the cache situation under different matrix sizes and optimizing the function step by step. Especially in the optimization step for 64*64 matrix, a little more complex but Intuitive approach was figured out and meet the requirements successfully. In the optimization for 61*67 matrix, I found the best parameter for the block matrix by several testing. Besides, all the optimization steps and my thoughts are illustrated in detail with necessary figures in the analysis part.
- To take care of the readability of the codes, I added detailed comments in each critical step.
- In brief, the project helped me have a better understanding of the cache performance. It provided me with an opportunity to apply theoretical knowledge to practice, which both enhanced my grasp of the knowledge and strengthened my coding ability. By the way, completing the whole project on my own gives me a great sense of achievement.

Finally, I would like to appreciate Miss Shen and teaching assistants for their careful guidance and support, from which I have benefited a lot.